

FCAI


fcai.fi

Model Averaging

3

8

$$\pi^{Greedy} = \arg \max_{\pi} \mathbb{E}_{p(f|\mathcal{D})} \left[J(\pi; f) \right]$$







Deisenroth et al. (2011). PILCO: A Model-Based and Data-Efficient Approach to Policy Search. ICMML.

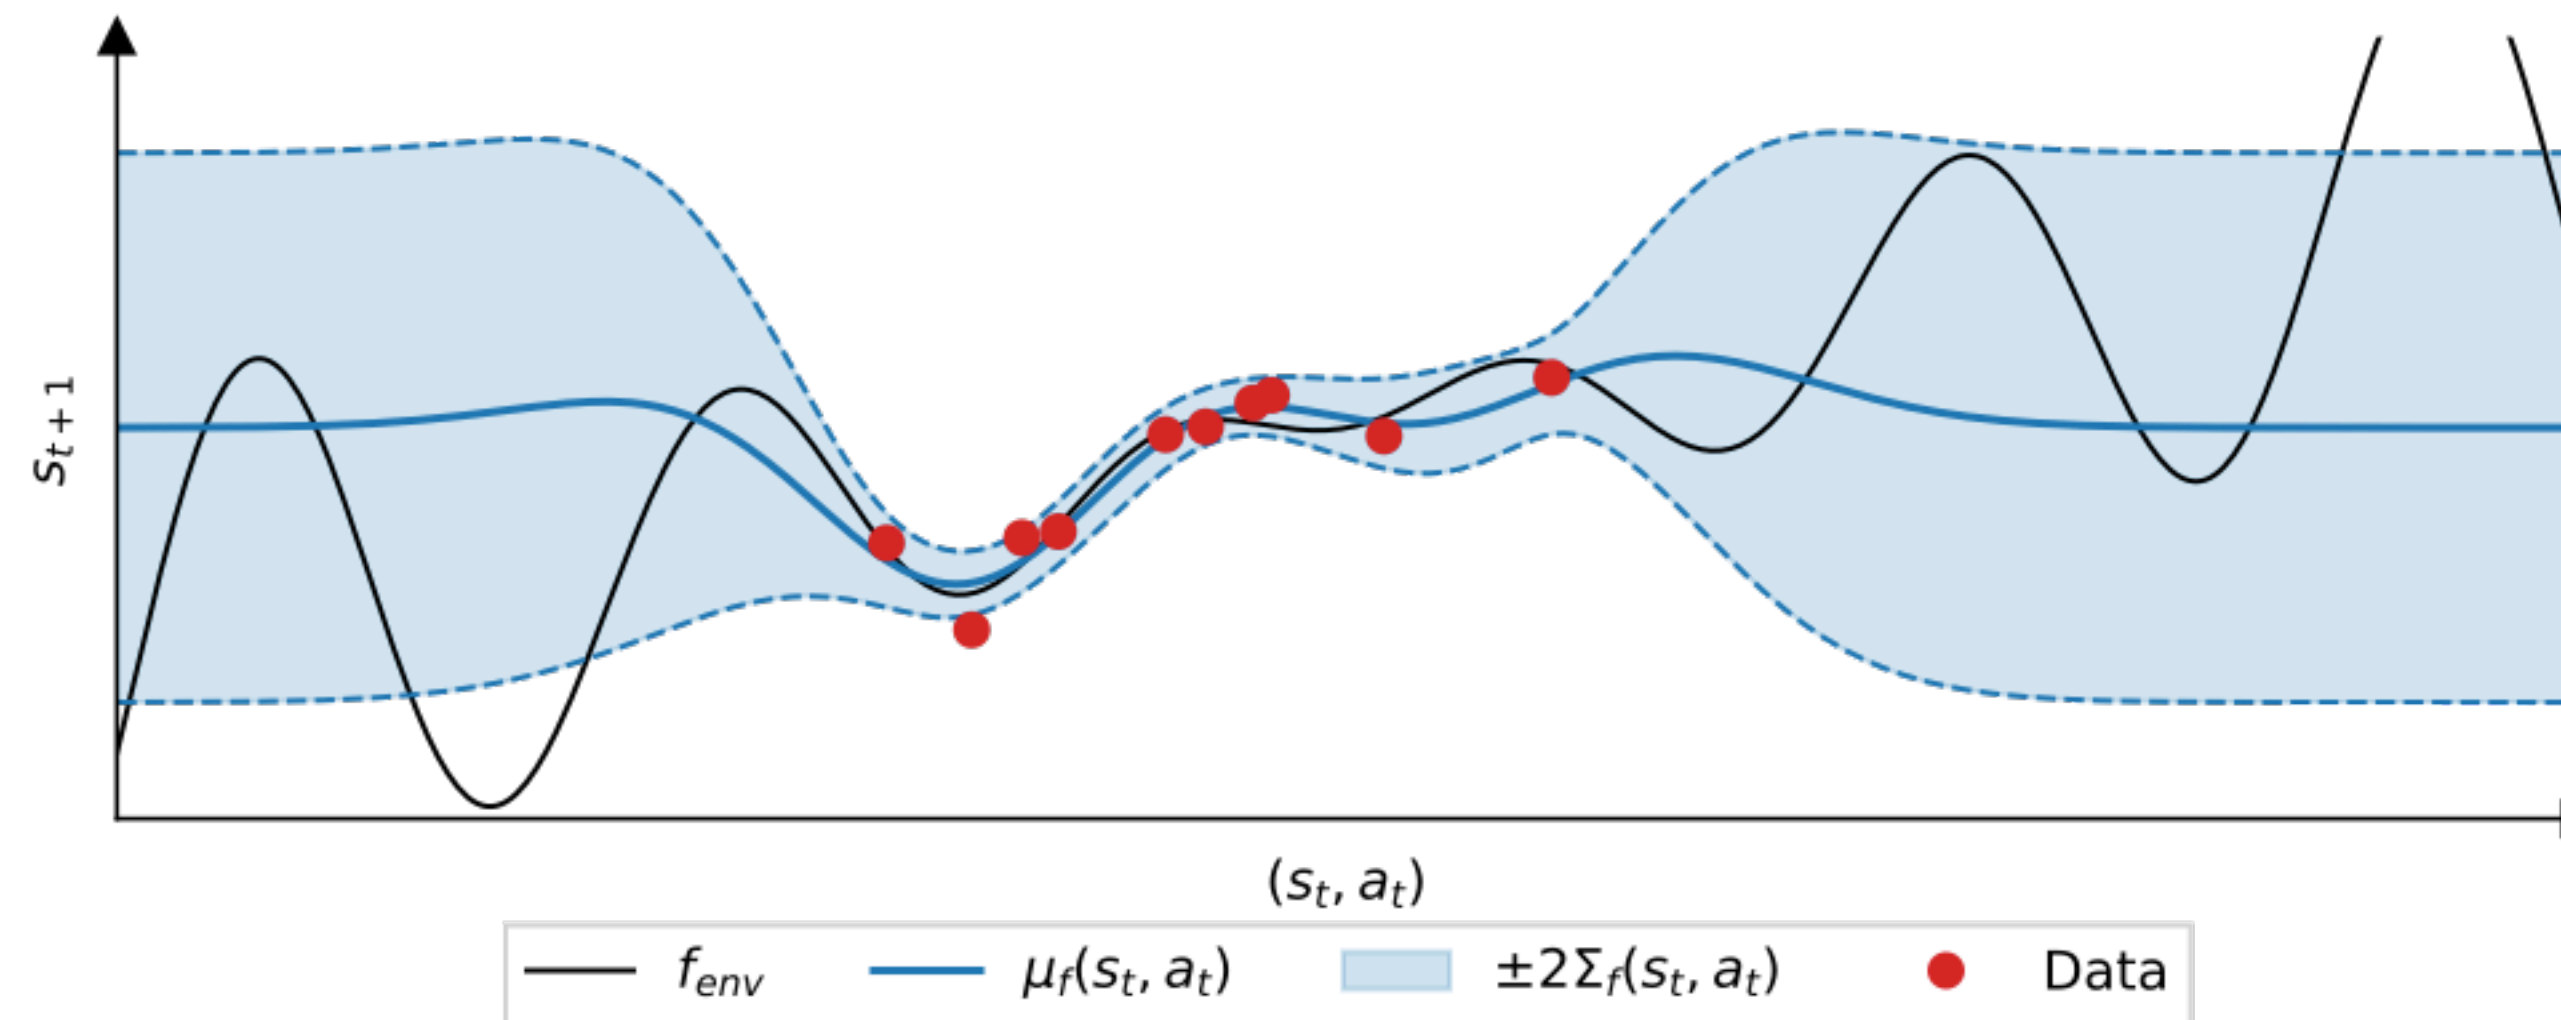
Kurtland et al. (2018). Deep Reinforcement Learning in a Handful of Trials using Probabilistic Dynamics Models. NeurIPS.

Model Averaging

$$\pi^{Greedy} = \arg \max_{\pi} \mathbb{E}_{p(f|\mathcal{D})} \left[J(\pi; f) \right]$$

PILCO, PETS, etc

Combats model bias



Exploration via Posterior Sampling

$$\pi^{PS} = \arg \max_{\pi} J(\pi; \tilde{f}), \quad \tilde{f} \sim p(f \mid \mathcal{D})$$

