2

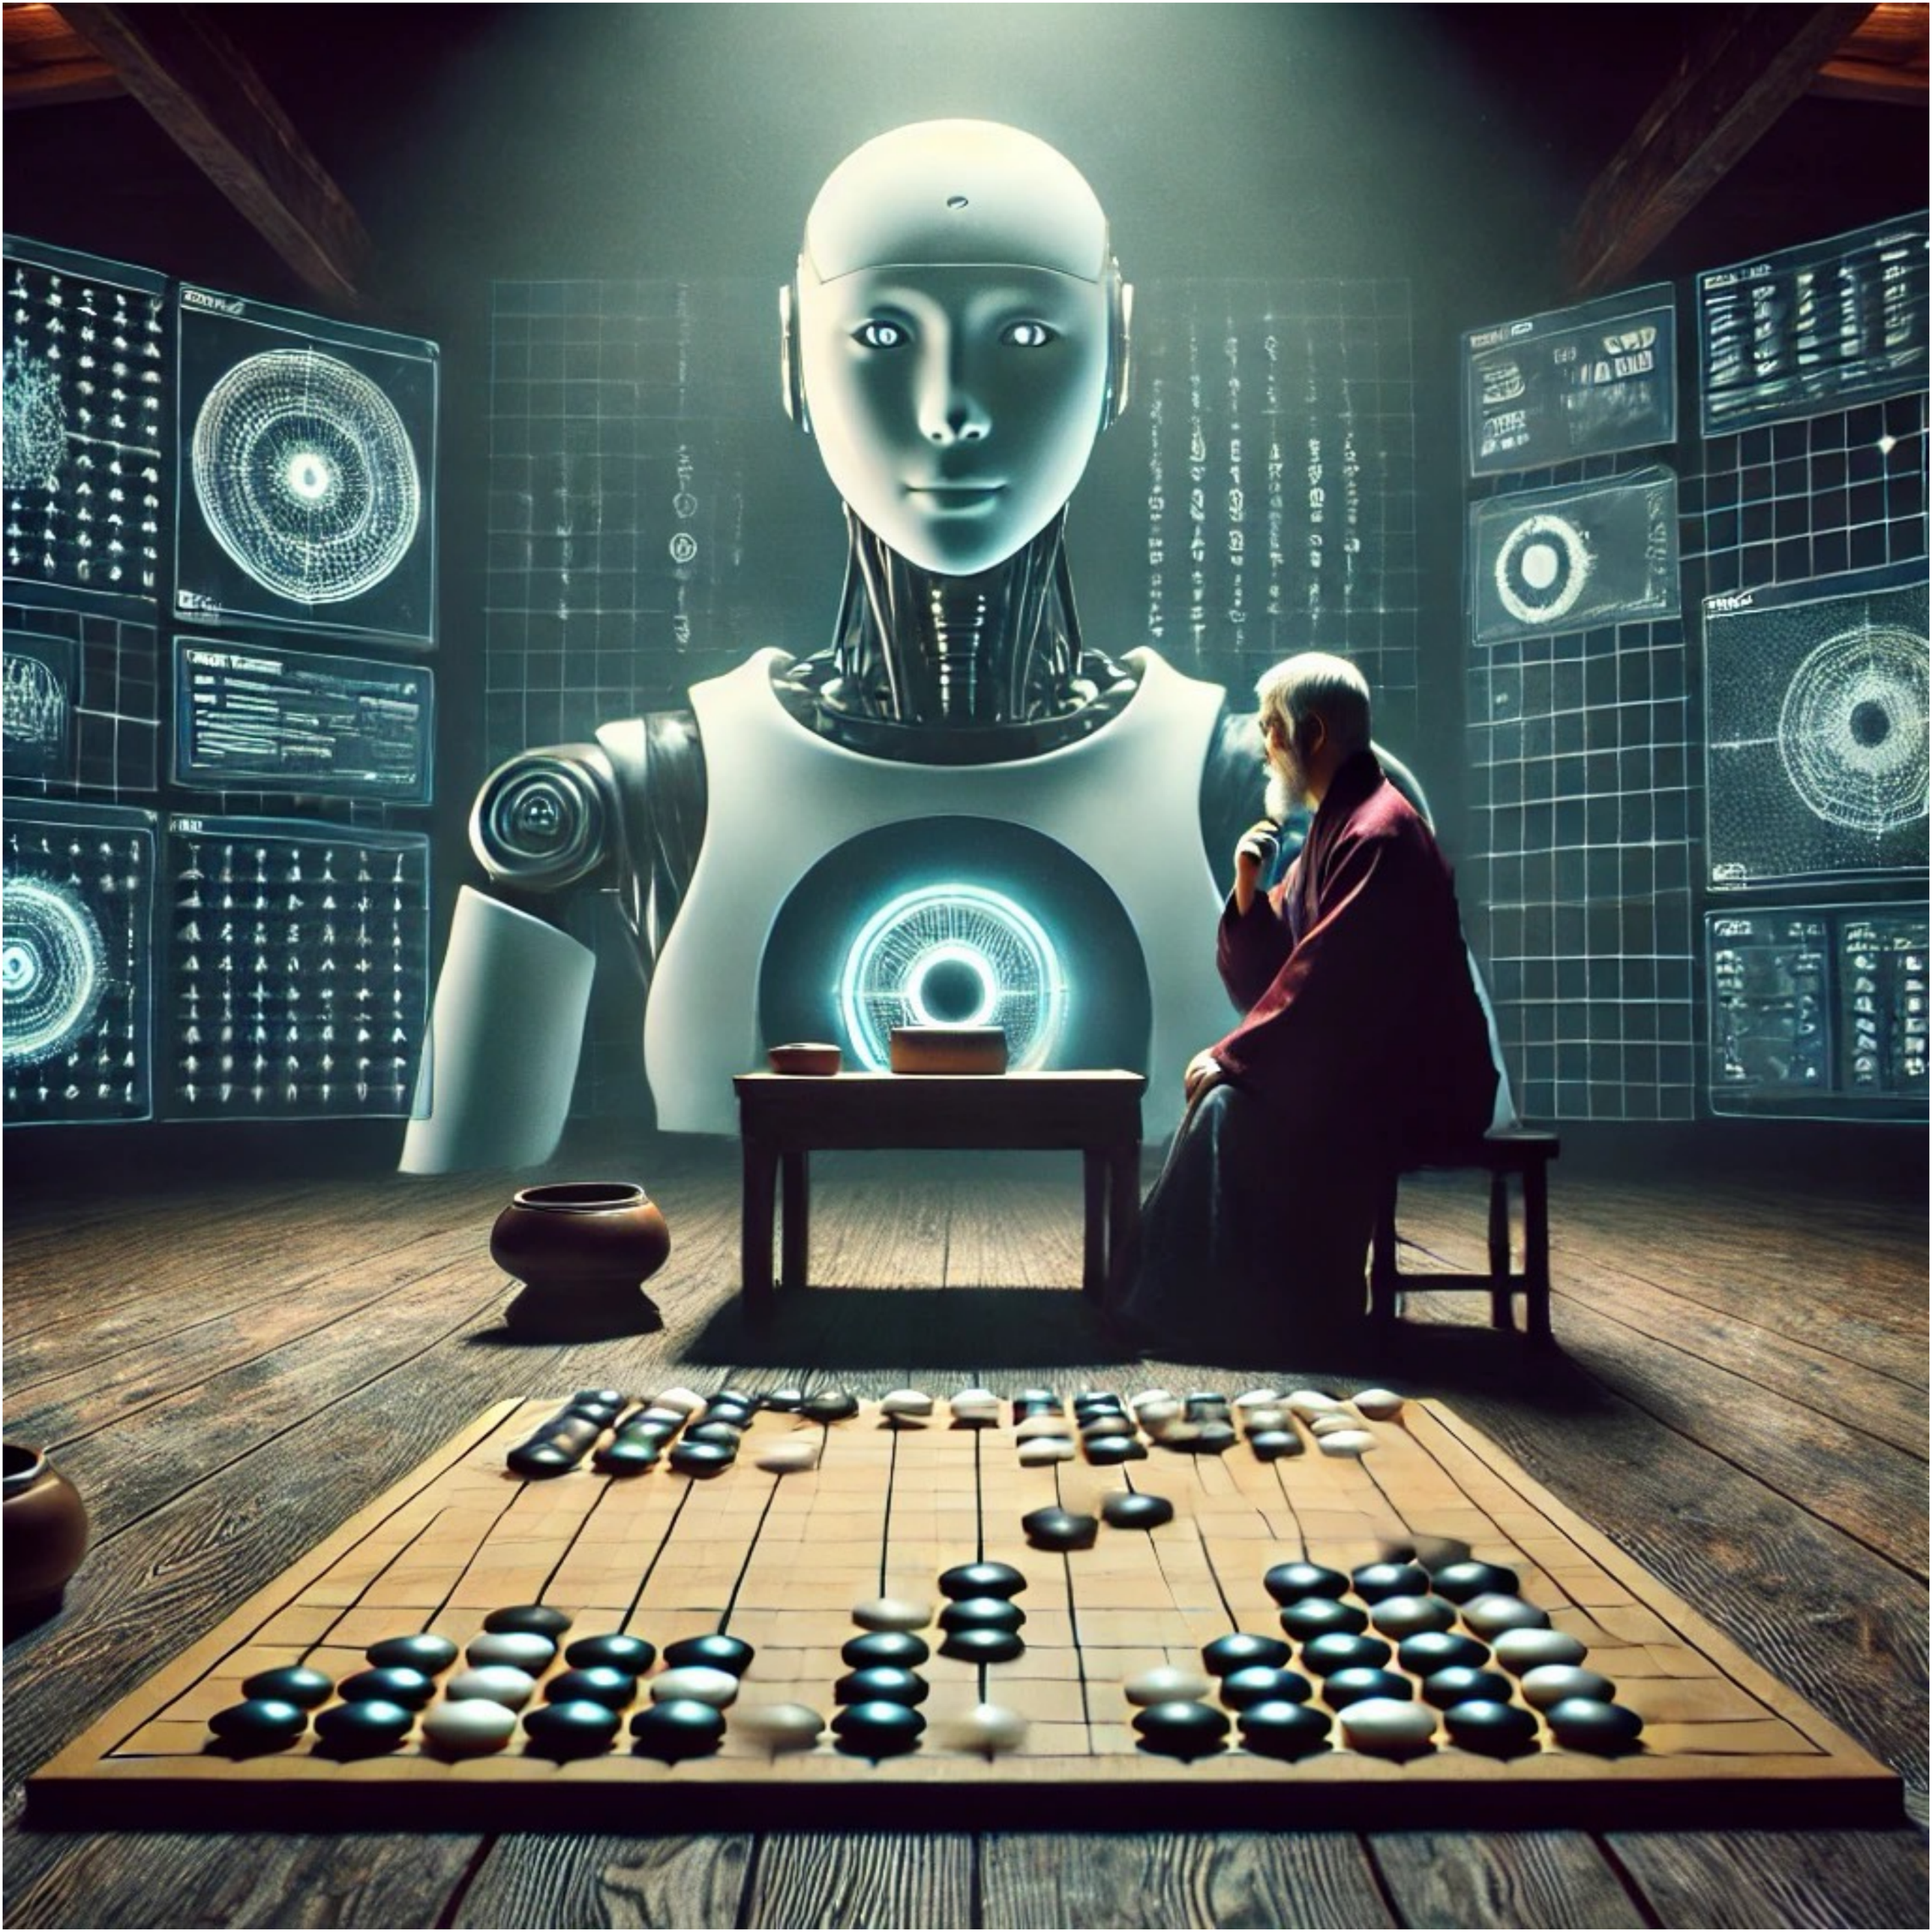

AlphaGo

Model-based reasoning for games

Silver et al. (2016). Mastering the game of Go with deep neural networks and tree search. Nature, 529(7587), 484.

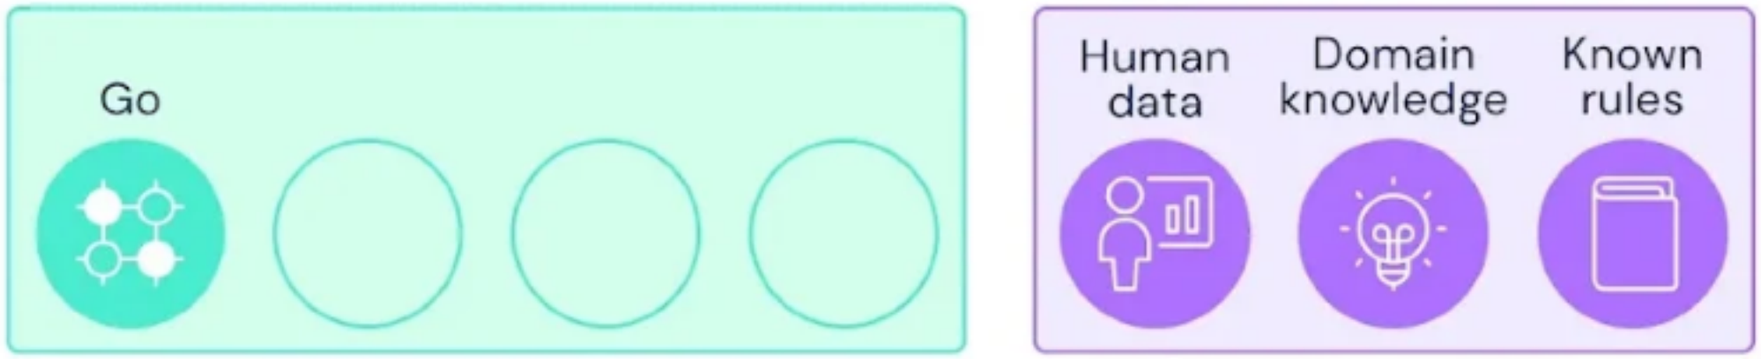

AlphaGo becomes the first program to master Go using neural networks and tree search
(Jan 2016, Nature)

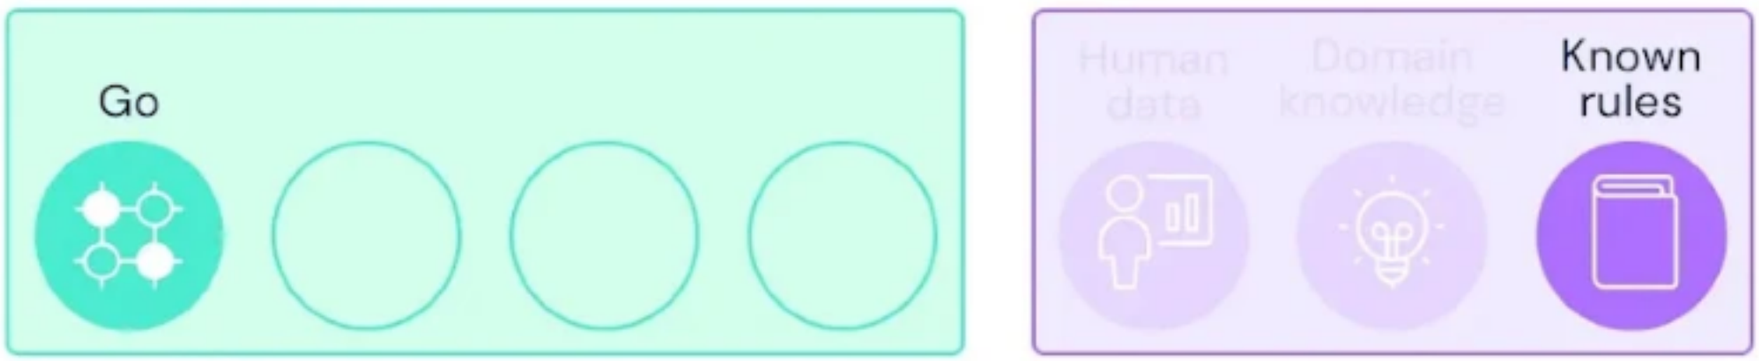

AlphaGo Zero learns to play completely on its own, without human knowledge
(Oct 2017, Nature)

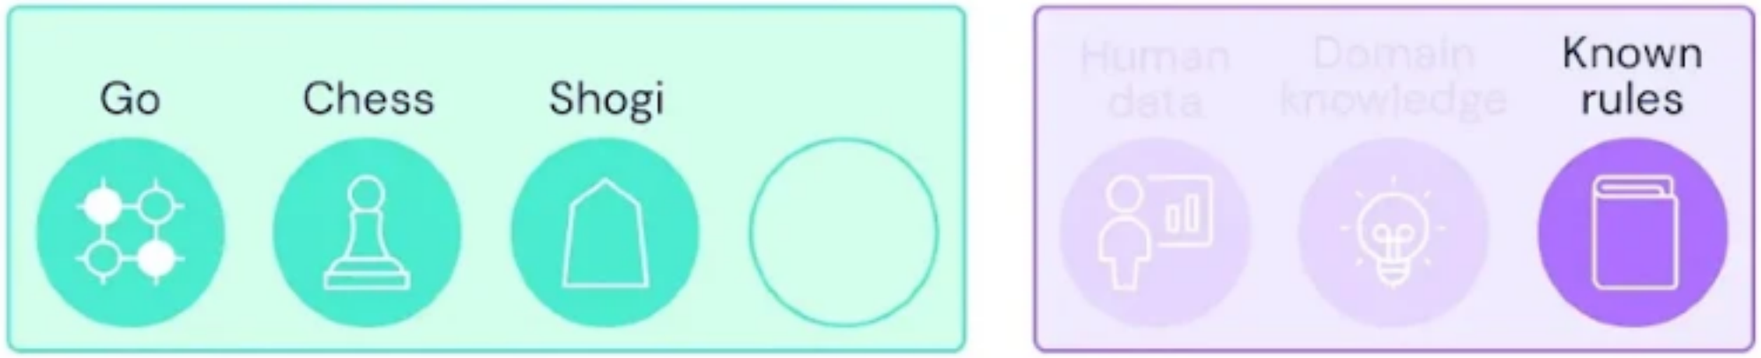

AlphaZero masters three perfect information games using a single algorithm for all games
(Dec 2018, Science)

MuZero learns the rules of the game, allowing it to also master environments with unknown dynamics.
(Dec 2020, Nature)

AlphaGo

Model-based reasoning for games

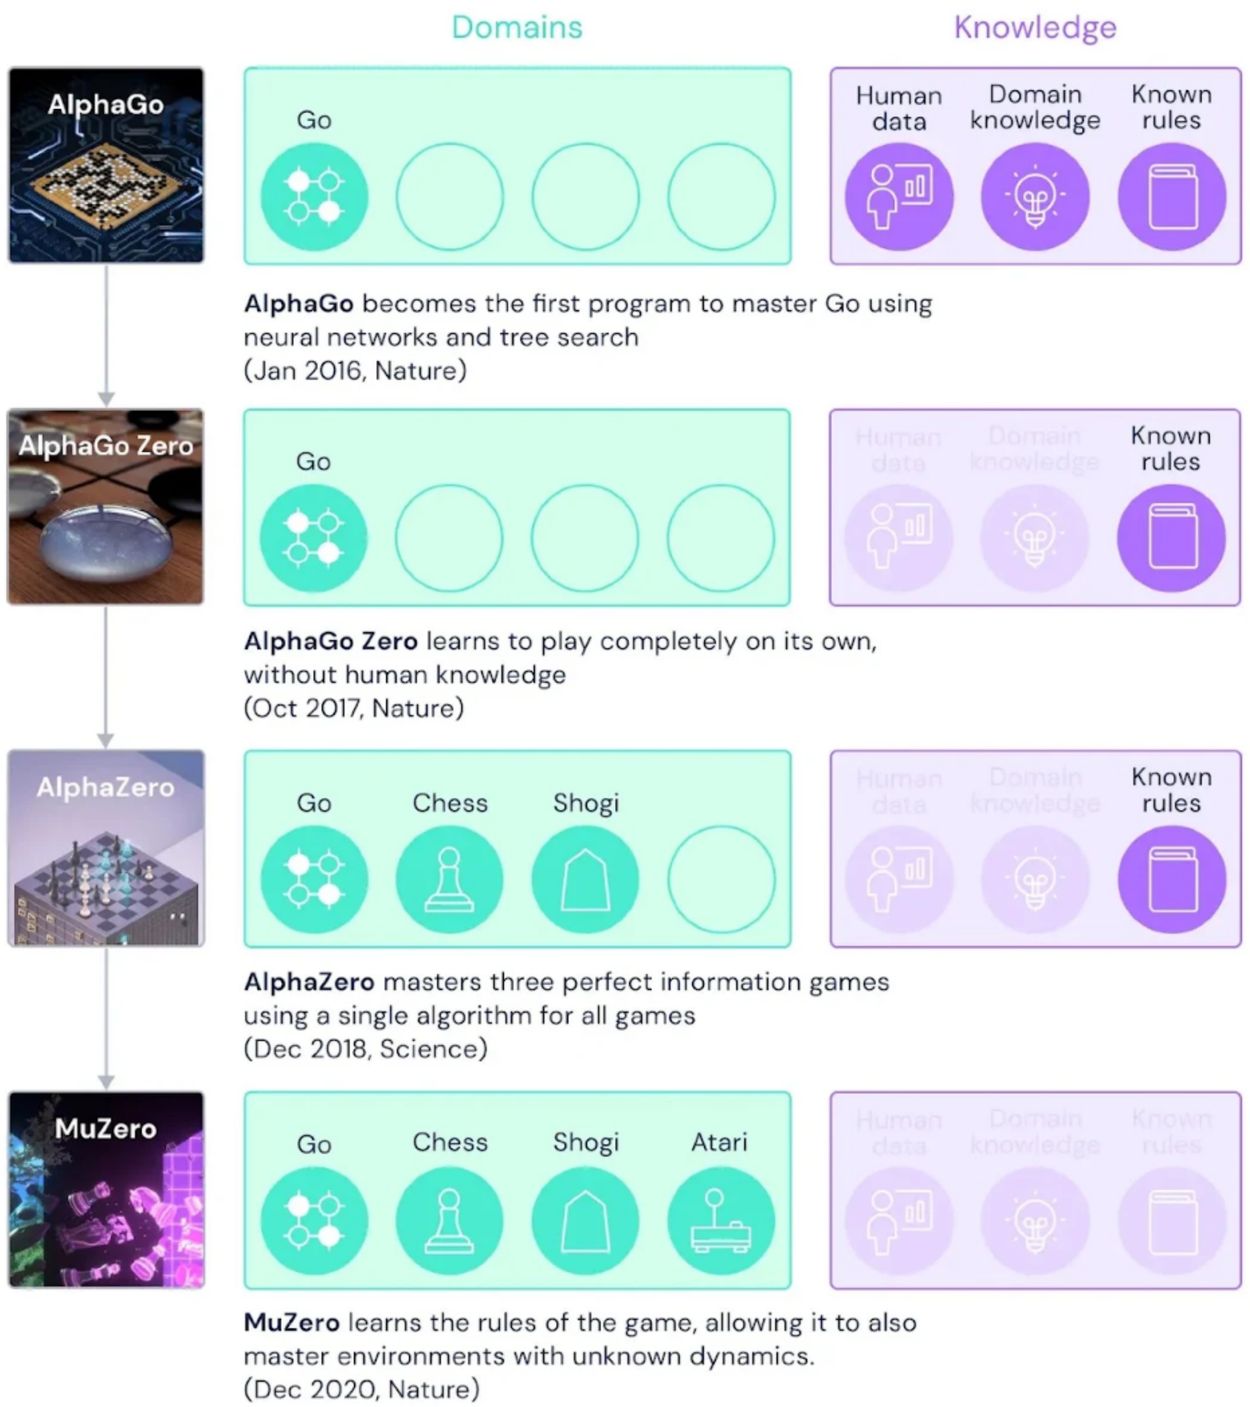

Silver et al. (2016). Mastering the game of Go with deep neural networks and tree search. Nature, 529(7587), 484.

FCAI

Machine Learning for Robotics

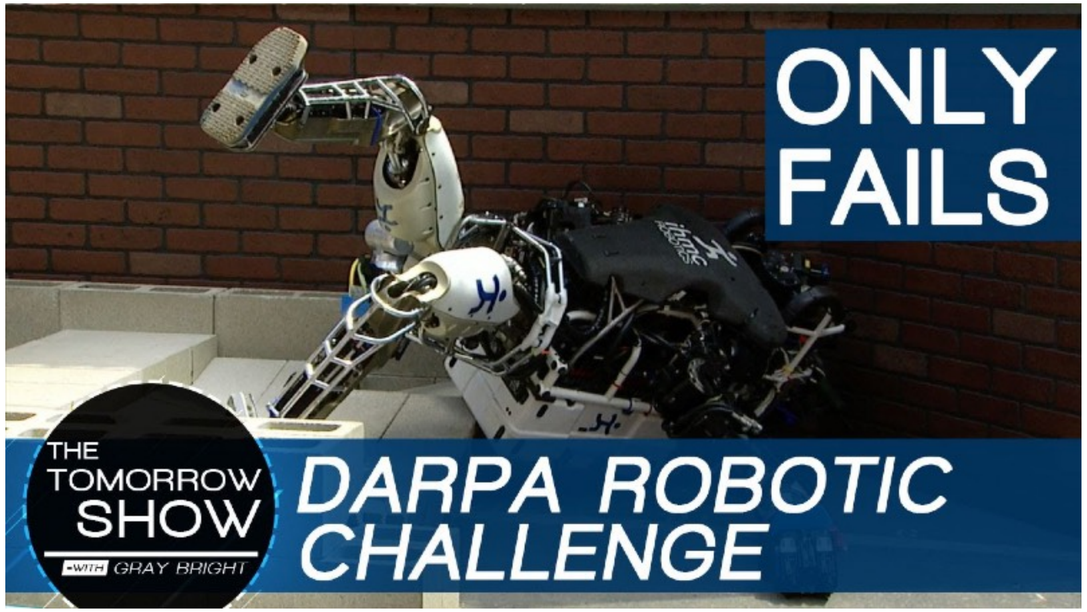

DARPA Robotics Challenge 2015

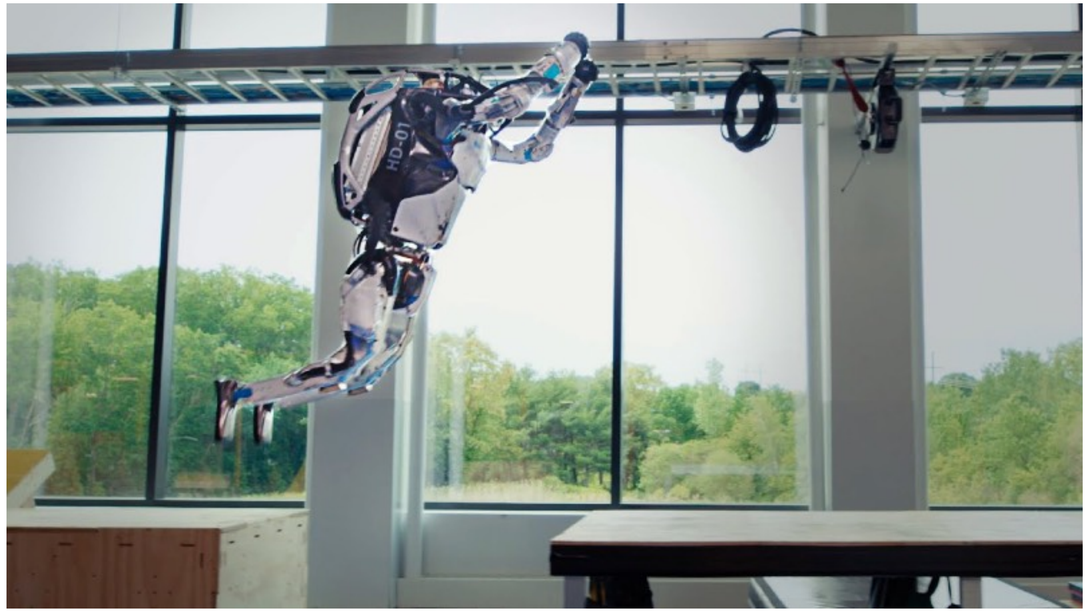

Boston Dynamics Atlas - Partners in Parkour